

WEEKLY PROGRESS UPDATE #4

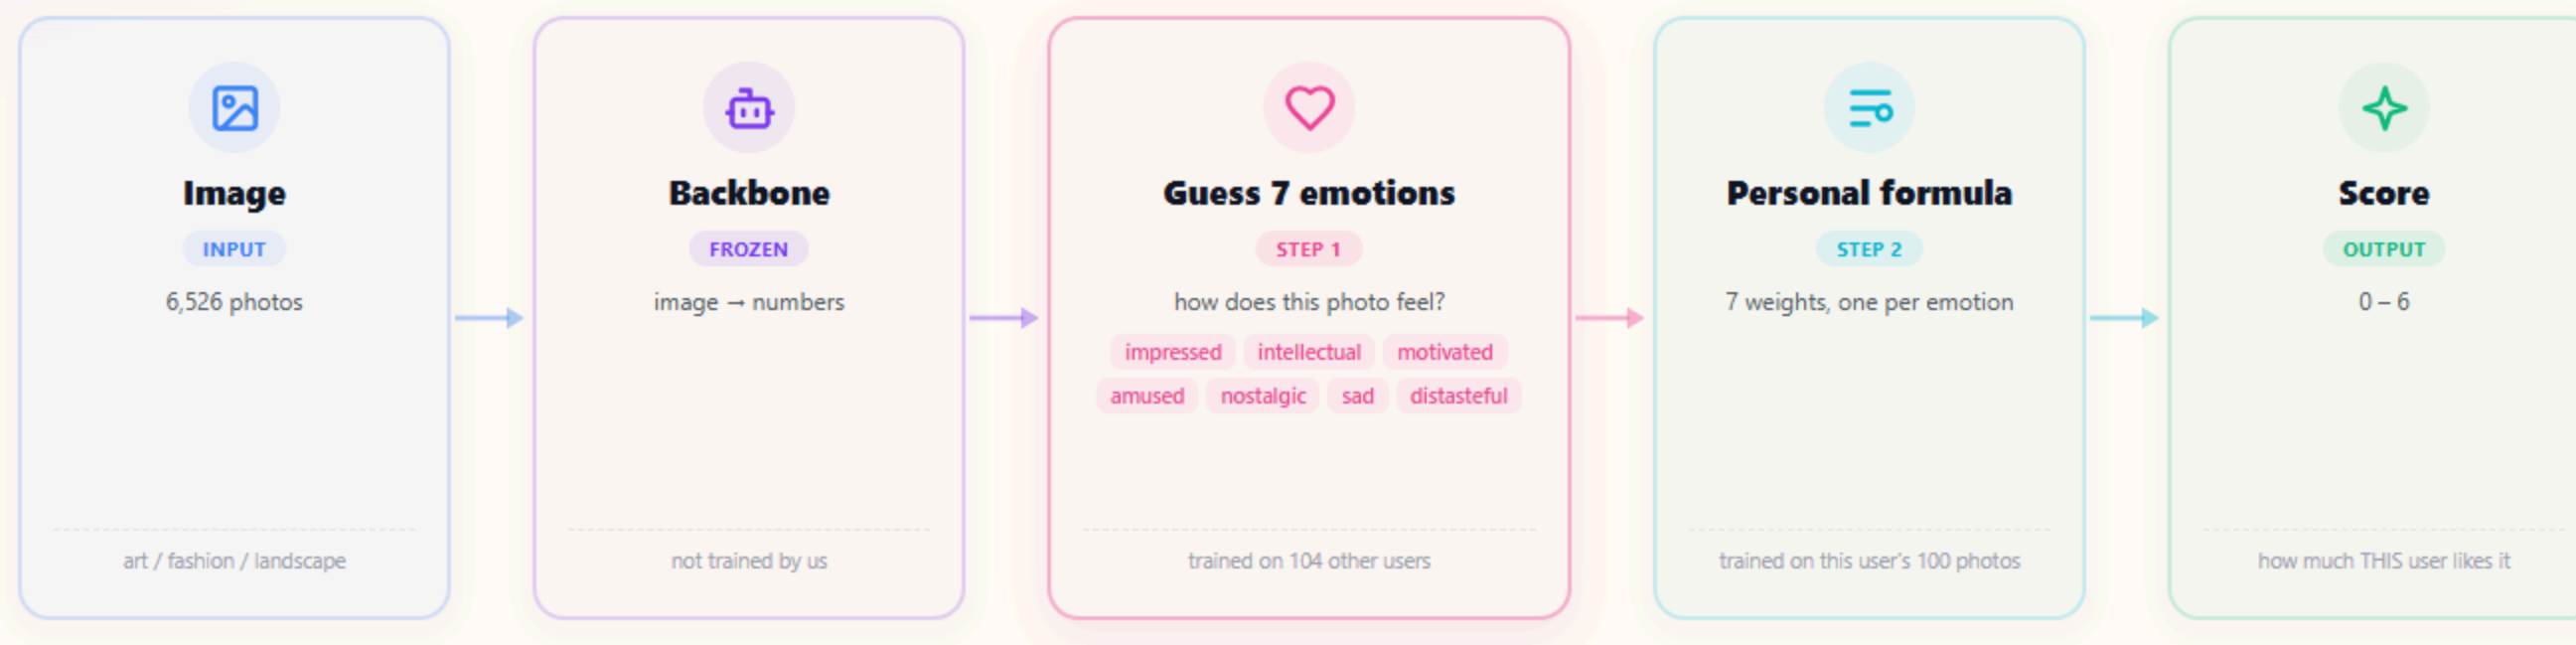
Emotion-Mediated Personalized Image Aesthetic Assessment

Predicting individual aesthetic preferences through emotional responses

How the system predicts a score.

Direct photo → score (skips the emotion step)

Hybrid photo → 7 emotions → score



Two separate steps, each fitted on different people

The emotion step learns from **104 other users**, so it is the same for everyone. The formula is fitted on the **target user's own 100 photos** — that is where personalisation happens. Both steps are simple ridge regression, not neural networks.

Terms & data

CCC — how close predictions are to real scores (0–1). Punishes wrong order *and* wrong size.

emo_r — how well we guess this person's emotions. Average now **0.27**.

XPASS-Vis — 129 users · 6,526 photos · 3 domains · 87,836 ratings · scores 0–6

Previous Finding & Today's Goals.

21 July 2026

LAST TIME (UPDATE #3)

Baseline Performance & Qwen Findings

Qwen3-VL Finding (Ryu & Yanaka)

Tested 8B: The stronger the backbone, the smaller the Hybrid-Direct gap.

Strongest Overall Baseline

Fine-tuned CLIP remains our highest performing model (CCC $\approx$ 0.400).

TODAY (UPDATE #4)

01 Noise Ceiling Analysis

Uncovering the true upper bound by accounting for human self-agreement.

02 When emotion help

In which cases does emotion help users, and in which cases does it hurt?

03 Qwen3-VL-4B Model

Testing the smaller 4B model recommended in Ryu & Yanaka to check compute vs accuracy.

Is that ceiling **actually real?**

Our best model

0.400

Human self-agreement

0.693

P-oracle (inflated)


0.725

PEOPLE DO NOT AGREE WITH THEMSELVES

The dataset has a built-in test-retest design: users rated the **SAME image twice** under identical conditions. People gave the exact same score only **44%** of the time.

4,509 Repeated Pairs Analyzed

Honest Upper Bound: The P-oracle (0.725) sits **above** human self-agreement because emotions & scores share same-session state noise. The real ceiling is **~0.69** — so our 0.400 reaches **~58%** of what is actually reachable.



From an average to a reason.

USER LEVEL 129 users × 3 domains = 387 units

1 Which person does it help?

Checking if emotion helps all users equally. We run a paired Wilcoxon test to see if it is a broad, common effect or driven by a few outliers.

2 What decides whether it helps?

Analyzing key factors that determine whether emotion helps:

A Emotion prediction accuracy (**emo_r**)

B Baseline strength of the Direct model (**Direct**)

C Rating consistency & category headroom

IMAGE LEVEL per-image error reduction

3 Which images does it help?

Checking if photos triggering different emotions gain different benefits. We group images by their dominant emotion (e.g. distasteful vs. intellectual) to compare.

The setup.

WHAT WE ARE EXPLAINING

Delta (**delta**)

$\text{delta} = \text{Hybrid CCC} - \text{Direct CCC}$

+ **positive delta**
Emotion helped

- **negative delta**
Emotion hurt

129 users × 3 domains = **387 units** to analyze

WHAT MIGHT EXPLAIN IT

How well we read feelings (**emo_r**)

Measures how accurately Stage 1 predicts **THIS specific user's** 7 emotions:

$\text{emo_r} = \text{mean correlation}(\text{predicted}, \text{actual}) \text{ across } 7 \text{ emotions}$

Average across all users

0.27

the weak link in the whole system

It helps almost everyone.

WHAT THE HISTOGRAM SHOWS

Each bar is one user × domain unit. Bars to the **right of zero** = emotion helped. Bars to the **left** = emotion hurt.

91.7% of units sit right of zero

p-value
< 0.001

Median gain
+0.071

Effect r
0.82

THE LEFT TAIL (8.3%, 32 UNITS)


Not random. They cluster exactly where emotion prediction was poor:

Left-tail avg emo_r
0.153

VS

Rest avg emo_r
0.283

→ So emotion accuracy (emo_r) seems to decide whether it helps or hurts. Is that true?



Who does it help, and why?

OUR GUESS BEFORE TESTING

Emotions should help more when we guess that person's feelings well — and hurt when we guess them badly.

Across all 387 users
 $r = +0.30$ ($p = 1e-9$)



HARDEST TO READ (BOTTOM 25%)
emo_r $\approx$ 0.09

+0.040 gain
80% helped

EASIEST TO READ (TOP 25%)
emo_r $\approx$ 0.45

+0.092 gain
97% helped

The 32 users it made worse: for them we read feelings at only **0.15**, against **0.28** for everyone else. Emotions hurt exactly where we read them badly.



How good must the emotion guess be?

Is there an accuracy cut-off below which emotion isn't worth it?

No Threshold in Our Data

Even the group we read worst still gains **+0.031** on average. The fitted line only reaches zero at **emo_r = -0.29**, which is below anything we actually see in the data.

DOSE-RESPONSE MODEL

emo_r controls **HOW MUCH** it helps and **HOW RISKY** it is


Worst Quartile Risk

31% hurt

Best Quartile Risk

3% hurt

Conclusion: Higher emotion accuracy increases expected gain while minimizing risk of negative delta, but emotion mediation remains net positive across all observed emotion accuracy ranges.



In other words: reading feelings better does not switch the benefit on — it just makes it bigger, and safer.

Only one thing really explains it.

WHAT IT LOOKS LIKE AT FIRST

Users whose plain model is already good also seem to gain more from emotions ($r = +0.15$).

But those two go together: $r = +0.66$. If we read someone's photos well, we do well at both.

WHAT HAPPENS WHEN WE HOLD EMO_R FIXED

delta $\leftrightarrow$ Direct CCC | controlling for emo_r


-0.08 (gone)

delta $\leftrightarrow$ emo_r | controlling for Direct CCC

+0.26 (SURVIVES)

Single Mechanism: Direct strength was a statistical confound. There is only ONE underlying mechanism controlling emotion benefit: **emotion predictability (emo_r)**.

Direct-CCC was a confound: its apparent effect ran through emo_r



Analogy: Foot size correlates with reading ability in children — both grow with age. Here: age = **emo_r**, foot size = **Direct CCC**, reading = **delta**.

Why is art different?

EMO_R → DELTA GAIN CORRELATION BY DOMAIN

Fashion $r = +0.38$ ✓ (significant)

Landscape $r = +0.34$ ✓ (significant)


Art $r = +0.14$ ✗ (not significant, $p = 0.11$)

H1: CEILING EFFECT

Art Direct baseline is already high (**0.384** vs fashion **0.202**). Less headroom → delta varies less → correlation harder to detect statistically.

H2: DIFFERENT MECHANISM

Aesthetic judgments in artistic works stem from non-emotional features like technique, composition, or historical context rather than emotions.



Testing the two possible reasons.

Three tests — do they point to "ceiling" or "different mechanism"?

1. Partial correlation — control for Direct-model strength

If it's a ceiling, art's correlation should recover once baseline strength is removed

art: +0.13 (ns) $\rightarrow$ +0.18 $\cdot$ $p = 0.048$

leans ceiling

2. Spread of the gain within each category (SD)

If art is capped, its gain should vary the least — lowest spread

art SD = 0.054 (lowest) $\cdot$ Levene $p = 0.21$

weak / borderline

3. Interaction test — is the emo_r $\rightarrow$ gain slope different in art?

If art had a different mechanism, this would show up clearly

$p = 0.056$ (not significant)

no evidence for H2

All three lean toward "ceiling" — but every number sits on the 0.05 border

So: no evidence art works differently — not proof they're the same

Which photos benefit most?

GROUPED BY DOMINANT EVOKED EMOTION


Distasteful	+0.043 (helps most)
Impressed	+0.021
Amused	+0.014
Nostalgic	+0.006
Intellectual	+0.001 (helps least)

STABILITY HYPOTHESIS & POWER LIMIT

Hypothesis: Images evoking stable/consistent emotions benefit more.

Spearman $r = +0.39$, $n = 7$	$p = 0.38$ (INCONCLUSIVE)
--------------------------------	---------------------------

Honest Reporting: With only $n=7$ emotion categories, the statistical test lacks sufficient power. Direction matches hypothesis (+0.39), but result must be reported as inconclusive.



Does the small model keep up?

MODEL VALIDATION & LAYER SWEEPING


Ryu & Yanaka noted 4B performed best on AADB. We evaluated whether this holds on XPASS-Vis.

4B Best layer we found: LT17 | 8B Best layer we found: LT15

PERFORMANCE COMPARISON (CCC)

Model (Dimensions)	Direct CCC	Hybrid CCC
Qwen4B (2560-d)	0.370	0.385 ⭐
Qwen8B (4096-d)	0.376	0.382

Efficiency Gain: The smaller 4B model (2,560 features) achieves slightly higher Hybrid CCC than 8B (4,096 features), saving significant GPU compute while maintaining performance.




The better the model, the less emotions add.

CORE FINDING ACROSS ALL TESTED BACKBONES

The stronger the vision backbone, the less explicit emotion-mediation adds on top.

Linear Correlation
r = -0.95
across 6 independent models



Implicit vs Explicit Representations

Strong VLMs already encode emotional semantics implicitly within their feature space. Making emotions explicit in Stage 2 yields diminishing returns as backbone capacity grows.

Multi-Backbone Rigor: This systematic trend only emerges when evaluating multiple diverse backbones (CLIP, ViT, ResNet, Qwen8B, Qwen4B) rather than relying on a single architecture.

To sum up.

- 01 The true human ceiling is about 0.69, not 0.725**

Measuring 4,509 image rating pairs where the same person rated the exact same photo twice shows a self-agreement of 0.693. The old 0.725 ceiling was inflated by shared mood noise because emotions and scores were rated in the same sitting.
- 02 Emotion helps 92 percent of people**

Expected gain is a gradual effect driven by how well we read their feelings, not a fixed cutoff, and not about how strong the plain model is on its own. Even the hardest-to-read group still sees a small average gain (+0.031).
- 03 Direct model strength is not the real driver of the gain**

Direct strength correlates with gains at $r = +0.15$, but it moves closely with emo_r at +0.66. Removing emo_r 's influence via partial correlation drops Direct's correlation to -0.08 (non-significant), showing that only emotion accuracy truly matters.
- 04 Art difference is a ceiling effect, not a new mechanism**

Art looks different mostly because it's already close to its ceiling, not because it works by a different mechanism—though the evidence there is borderline. Controlling for baseline strength via partial correlation restores significance (+0.175, $p=0.048$).
- 05 Stronger backbones need emotion less**

A near-perfect negative correlation ($r = -0.95$) across 6 models confirms that powerful backbones already capture emotional features. The smaller 4B model holds up fine, adding a sixth data point to this clean pattern.

What comes next.

WHAT I PLAN TO TRY NEXT

1

Cold-start analysis

How many ratings needed before personal model beats population formula?

2

Trait-conditioned emotion prediction

Can trait vectors predict how a user's emotions deviate from population? If yes → zero-shot personalization.

Not in a hurry — can continue from Thailand after the internship if needed.

What I would like **your view on.**

- 1 Should we run cross-dataset validation?**
Should I validate the Hybrid pipeline on other datasets (like FLICKR-AES or PARA) to verify generalization, and do you have recommendations? Or is evaluating on a single unique dataset generally accepted in this community?
- 2 Should we fine-tune Qwen-4B end-to-end?**
Is end-to-end fine-tuning on the recommended Qwen-4B model worth the computation for potential accuracy gains, or should we freeze the model here and start writing the paper?
- 3 Does the ICI baseline actually benefit from traits?**
Should we run an ablation study by removing personality traits from the ICI baseline to verify whether its accuracy gain comes from traits or its architecture?
- 4 What is our weakest point for an ACM MM submission?**
If we submit this work to ACM MM 2027, what do you think is our main weakness or the most likely reason reviewers might reject it?
- 5 ACM MM vs. ACII: Which is more suitable?**
Which venue fits this project's profile better: a major multimedia conference like ACM MM, or an affective computing focus like ACII?

PROGRESS UPDATE #4

Thank You

Pinwa · JAIST Internship

Backup: how the analysis is set up.

QA-1 How is delta computed?

A per-user pairwise difference, taken for every user in every domain from results that already existed.

$$\Delta_{u,d} = CCC_{Hybrid\,u,d} - CCC_{Direct\,u,d}$$

Data source

exp1_per_user

Rows

387

129 users × 3 domains

Retraining

Zero

secondary analysis

Strict controls: Hybrid and Direct come from the exact same test split, the same user ratings and identical folds.

QA-3 Is there any data leakage?

The pipeline is fitted in two stages, on two groups of users that never overlap.

Stage 1 • General set, 104 users

Trained exclusively on general users to learn the generic image
→ emotion mapping.

Stage 2 • Target set, 25 users

Personal ridge model fit strictly per target user, on their own
training fold only.

Zero user overlap: Leakage would require seeing target-user ratings during Stage 1, which this split makes impossible.

QA-2 Isn't Direct CCC a second cause?

A backbone that represents images well is good at **both** direct scoring and emotion prediction, so the two move together ($r = +0.66$).

Uncontrolled correlation

+0.15

looks like a second factor

Partial, controlling emo_r

-0.04

effect vanishes

Classic confound: Like shoe size correlating with reading ability when age is the real cause. Direct CCC tracks delta only because both are driven by emo_r.

QA-4 What exactly is emo_r?

Emotion predictability for one specific user: how well Stage 1 recovers their 7 emotion ratings.

$$\text{emo_r}(u) = \frac{1}{7} \sum_{e=1}^7 r(\hat{e}_e, e_e)$$

Stage 1 bottleneck

0.27

mean emotion accuracy

Stage 2 headroom

0.693

P-oracle, true emotions

The bottleneck: Stage 2 could reach 0.69 given perfect emotions, so the pipeline is capped by Stage 1.

Backup: statistics and metrics.

QA-5 **p = 0.056 — proven or not?**

$\Delta \sim \text{emo_r} \times \text{domain} \rightarrow p = 0.056$ (not significant)

Formal decision

We fail to reject the null at alpha = 0.05. We do **not** have statistical evidence that the mechanism in art differs from fashion or landscape.

Scientific nuance

Failing to reject is **not** proof the mechanisms are identical. It sits on the border, so we state it as inconclusive, favouring ceiling.

Methodological rigor: Absence of evidence is not evidence of absence. We report the value transparently.

QA-7 **Which cases failed?**

Characterising the 32 user × domain units where emotion mediation made the prediction worse.

Failure proportion

32 / 387

8.3% of units

Mean emo_r, hurt

0.153

read badly

Mean emo_r, helped

0.283

read well

Root cause: Mediation fails almost exclusively where Stage 1 misses the user's emotion profile, which injects noise into Stage 2.

QA-6 **Did you correct for multiple comparisons?**

28 candidate predictors of delta were screened, so the significance threshold has to be tightened.

$$\alpha_{\text{adj}} = \frac{0.05}{28} = 0.00179$$

Predictors

28

traits, stats, baselines

Method

Bonferroni

Survivor

emo_r

p = 1e-9

Robust significance: Only emotion-reading accuracy survives strict correction, accounting for $R^2 = 0.13$ of delta variance.

QA-8 **Why CCC and not RMSE or Spearman?**

Three candidate metrics, and what each one is blind to.

Lin's CCC

Ranking agreement **and** absolute scale. Punishes the right order at the wrong scale.

RMSE / MSE

Absolute error only. Does not isolate linear agreement or ranking quality.

Spearman rank

Ordinal ranking only, ignores scale entirely. Reported as a secondary metric.

Literature standard: CCC is the benchmark used by Hayashi-san et al. (XPASS-Vis) and Ryu & Yanaka (2024).